



Wall-clock time by runtime (lower is better, log scale; black outline = fastest)


Word tokenize
(1 MB)




Sentence split
(1 MB)




Porter stem
(100k words)




Bigram PMI
(1 MB)




FreqDist count
(1 MB)




Naive Bayes
train+eval



Everygrams n=1-3
(200 sents)



KN-3 LM
perplexity



Python NLTK 3.10 bun_nltk native bun_nltk WASM Node.js + WASM npm natural